

**General constraints for code submissions** Please adhere to these rules to make our and your life easier! We will deduct points if your solution does not fulfill the following:

- If not stated otherwise, we will use exclusively Python 3.5.
- If not stated otherwise, we expect a Python script, which we will invoke exactly as stated on the exercise sheet.
- Your solution exactly returns the required output (neither less nor more) – you can implement a `--verbose` option to increase the verbosity level for developing.
- Add comments and docstrings, so we can understand your solution.
- (If applicable) The **README** describes how to install requirements or provides addition information.
- (If applicable) Add required additional packages to **requirements.txt**. Explain in your **README** what this package does, why you use that package and provide a link to it's documentation or GitHub page.
- (If applicable) All prepared unittests have to pass.
- (If applicable) You can (and sometimes have to) reuse code from previous exercises.

---

Now that you have learned about the application of reinforcement learning for learning to learn and algorithm control, you will implement a basic RL agent yourself.

**1. Q-Learning** [9 points]

We present you a simple environment (based on the OpenAI gym standard), in which your agent has to learn to approximate a Sigmoid function.

(a) Implement a tabular function approximator that is able to handle float data, by mapping float state feature to the closest integer. [2pt.]

(b) Implement the Q-learning algorithm. Train your agent on the provided environment in its default configuration for 10 000 training episodes. Your agent will have to learn a policy of length 11 with action space  $\{0, 1, 2\}$ . The optimal policy will result in a reward of roughly 10.16 and looks like 0, 0, 0, 0, 1, 1, 1, 2, 2, 2, 2. [5pt.]

If your agent at first does not produce the optimal policy play around with the learning rate  $\alpha$  and the exploration factor  $\epsilon$ . In your PDF report all learning rates and exploration factors you tried and the resulting policies.

In your own words explain how the learning rate and the exploration factor influence the Q-learning on this deterministic environment.

(c) In the same folder you called **main.py** from, the results of training will be stored. Plot the true reward against the expected reward for all learning rates and exploration factors you tried and add them to your PDF. [2pt.]

**2. Feedback** [Bonus: 0.5 points]

For each question in this assignment, state:

- How long you worked on it.
- What you learned.
- Anything you would improve in this question if you were teaching the course.

**This assignment is due on 12.07.19 (10:00).** Submit your solution for the tasks by uploading a PDF to your groups BitBucket repository. The PDF has to include the name of the submitter(s).